

CKS Exam Questions and Answers

1. Mirror scan ImagePolicyWebhook
2. sysdig detects pods
3. clusterrole
4. AppArmor
5. PodSecurityPolicy
6. Network Policy
7. dockerfile detection and yaml file problem
8. pod security
9. Create SA
10. trivy detects mirror security
11. Create secret
12. kube-bench
13. gVsior
14. Audit
15. Default Network Policy
16. Falco detection output log format

kubernetes exam in action

Copyright statement: This article is an original article by CSDN blogger "ghostwritten" and follows the CC 4.0 BY-SA copyright agreement. Please attach the original source link and this statement for reprinting.

Original link: <https://blog.csdn.net/xixihahalelehehe/article/details/122525427>

Exam Information

2 hours

15-20 Topics

The appointment time is the same as CKA, and the results will be released within 32 hours

Out of 100, 87 or 93, but a passing score of 67

Simulation environment

4 environments 1 console

NAT network segment 192.168.26.0

Mock exam questions

1. Mirror scan ImagePolicyWebhook

switch cluster kubectl config use-context k8s

context

A container image scanner is set up on the cluster, but It's not yet fully integrated into the cluster's configuration When complete, the container image scanner shall scall scan for and

reject the use of vulnerable images.

task:

You have to complete the entire task on the cluster's master node, where all services and files have been prepared and placed

Given an incomplete configuration in directory /etc/kubernetes/aa and a functional container image scanner with HTTPS endpoint http://192.168.26.60:1323/image_policy

1.enable the necessary plugins to create an image policy

2.validate the control configuration and change it to an implicit deny

3. Edit the configuration to point the provided HTTPS endpoint correctly

Finally, test if the configuration is working by trying to deploy the vulnerable resource

/csk/1/web1.yaml

Problem solving ideas

ImagePolicyWebhook

Keywords: image_policy, deny

1. Switch the cluster, view the master, sshmaster

2. ls /etc/kubernetes/xxx

3. vi /etc/kubernetes/xxx/xxx.yaml change true to false

The address of https in vi /etc/kubernetes/xxx/xxx.yaml

volume needs to be mounted

4. Enable ImagePolicyWebhook and --admission-control-config-file=

5. systemctl restart kubelet

6. kubectl run pod1 --image=nginx

Case:

configure /etc/kubernetes/manifests/kube-apiserver.yaml

Add ImagePolicyWebhook related policies

Restart api-server, systemctl restart kubelet

Failed to verify image creation pod

Modify /etc/kubernetes/admission/admission_config.yaml policy defaultAllow: true

Revalidate the image to create the pod

\$ ls /etc/kubernetes/aa/

admission_config.yaml apiserver-client-cert.pem apiserver-client-key.pem external-cert.pem

external-key.pem kubeconf

1

2

\$ cd /etc/kubernetes/aa

\$ cat kubeconf

apiVersion: v1

kind: Config

clusters refers to the remote service.

clusters:

- cluster:

certificate-authority: /etc/kubernetes/aa/external-cert.pem # CA for verifying the remote

service.

server: http://192.168.26.60:1323/image_policy # URL of remote service to query. Must use 'https'.

name: image-checker

contexts:

- context:

cluster: image-checker

user: api-server

name: image-checker

current-context: image-checker

preferences: {}

users refers to the API server's webhook configuration.

users:

- name: api-server

user:

client-certificate: /etc/kubernetes/aa/apiserver-client-cert.pem # cert for the webhook admission controller to use

client-key: /etc/kubernetes/aa/apiserver-client-key.pem # key matching the cert

\$ cat admission_config.yaml

apiVersion: [apiserver.config.k8s.io/v1](https://kubernetes.io/docs/reference/generated/apiserver-api-types/)

kind: AdmissionConfiguration

plugins:

- name: ImagePolicyWebhook

configuration:

imagePolicy:

kubeConfigFile: /etc/kubernetes/aa/kubeconf

allowTTL: 50

denyTTL: 50

retryBackoff: 500

defaultAllow: false

#Modify api-server configuration

\$ cat /etc/kubernetes/manifests/kube-apiserver.yaml

.....

- command:

- kube-apiserver

- --admission-control-config-file=/etc/kubernetes/aa/admission_config.yaml #Add this line

- --advertise-address=192.168.211.40

- --allow-privileged=true

- --authorization-mode=Node,RBAC

- --client-ca-file=/etc/kubernetes/pki/ca.crt

- --enable-admission-plugins=NodeRestriction,ImagePolicyWebhook # #Modify this line

- --enable-bootstrap-token-auth=true

```

- --etcd-cafile=/etc/kubernetes/pki/etcd/ca.crt
.....
- mountPath: /etc/kubernetes/pki
  name: k8s-certs
  readOnly: true
- mountPath: /etc/kubernetes/aa #Add this line
  name: k8s-admission #Add this line
  readOnly: true #Add this line
.....
- hostPath: #Add this line
  path: /etc/kubernetes/aa #Add this line
  type: DirectoryOrCreate #Add this line
  name: k8s-admission #Add this line
- hostPath:
  path: /usr/local/share/ca-certificates
  type: DirectoryOrCreate
  name: usr-local-share-ca-certificates
- hostPath:
  path: /usr/share/ca-certificates
  type: DirectoryOrCreate
  name: usr-share-ca-certificates
status: {}

```

```

$ k get nodes
NAME STATUS ROLES AGE VERSION
master Ready control-plane, master 9d v1.20.1
node1 Ready <none> 9d v1.20.1
node2 Ready <none> 9d v1.20.1

```

```

# Failed to create pod
$ k run test --image=nginx
Error from server (Forbidden): pods "test" is forbidden: Post "https://external-service:1234/check-image?timeout=30s": dial tcp: lookup external-service on 8.8.8.8:53: no such host

```

```

#Modify admission_config.yaml configuration
$ vim /etc/kubernetes/aa/admission_config.yaml
apiVersion: apiserver.config.k8s.io/v1
kind: AdmissionConfiguration
plugins:
- name: ImagePolicyWebhook
  configuration:
    imagePolicy:
      kubeConfigFile: /etc/kubernetes/aa/kubeconf
      allowTTL: 50
      denyTTL: 50

```

```
retryBackoff: 500
defaultAllow: true #Modify this behavior true
```

```
#Restart api-server
$ ps -ef | grep api
root 78871 39023 0 20:17 pts/3 00:00:00 grep --color=auto api

$ mv ../kube-apiserver.yaml .

#Create pod successfully
$ k run test --image=nginx
pod/test created
```

2. sysdig detects pods

switch cluster kubectl config use-context k8s
you may use your browser to open one additional tab to access sysdig's documentation or
Falco's documentation

Task:

use runtime detection tools to detect anomalous processes spawning and executing frequently
in the single container belonging to Pod redis.

Two tools are available to use:

sysdig

falco

the tools are pre-installed on the cluster's worker node only; they are not available on the base
system or the master node.

using the tool of your choice (including any non pre-install tool) analyse the container's
behaviour for at least 30 seconds, using filters that detect newly spawning and executing
processes store an incident file at /opt/2/report, containing the detected incidents one per line in
the following format:

```
[timestamp],[uid],[processName]
```

```
1
```

Problem solving ideas

Sysdig User Guide

keyword: sysdig

0. Remember to use sysdig -l |grep to search for relevant fields

1. Switch the cluster, query the corresponding pod, and ssh to the node host corresponding to
the pod

2. Use sysdig, pay attention to the required format and time, and redirect the result to the
corresponding file

3. sysdig -M 30 -p "%evt.time,%user.uid,%[proc.name](#)" [container.id](#)=container id >/opt/2/report

case

3. clusterrole

switch cluster kubectl config use-context k8s

context

A Role bound to a pod's serviceAccount grants overly permissive permission

Complete the following tasks to reduce the set of permissions.

task

Given an existing Pod name web-pod running in the namespace monitoring Edit the Role bound to the Pod's serviceAccount sa-dev-1 to only allow performing list operations, only on resources of type Endpoints

create a new Role named role-2 in the namespaces monitoring which only allows performing update operations, only on resources of type persistentvolumeclaims.

create a new Rolebinding name role-2-binding binding the newly created Role to the Pod's serviceAccount

Problem solving ideas

RBAC

Keywords: role, rolebinding

1. Find the role modification permissions corresponding to rolebinding as list and endpoints

\$ kubectl edit role role-1 -n monitoring

2. Remember --verb is permission --resource is object

\$ kubectl create role role-2 --verb=update --resource=persistentvolumeclaims -n monitoring

3. Create a binding and bind it to the corresponding sa

\$ kubectl create rolebinding role-2-binding --role=role-2 --serviceaccount=monitoring:sa-dev-1 -n monitoring

4. AppArmor

switch cluster kubectl config use-context k8s

Context

AppArmor is enabled on the cluster's worker node. An AppArmor profile is prepared, but not enforced yet. You may use your browser to open one additional tab to access the AppArmor documentation. Task

On the cluster's worker node, enforce the prepared AppArmor profile located at /etc/apparmor.d/nginx_apparmor . Edit the prepared manifest file located at /cks/4/pod1.yaml to apply the AppArmor profile. Finally, apply the manifest file and create the pod specified in it

Problem solving ideas

apparmor

keyword: apparmor

1. Switch the cluster, remember to check the nodes, ssh to the node node
2. View the corresponding configuration file and name


```
$ cd /etc/apparmor.d
$ vi nginx_apparmor
$ apparmor_status |grep nginx-profile-3 # There is no grep to indicate that it is not started
$ apparmor_parser -q nginx_apparmor # Load and enable this configuration file
```
3. Modify the corresponding yaml to apply this rule, open the URL copy example of the official website, modify the container name and the local configuration name


```
$ vi /cks/4/pod1.yaml
apiVersion: v1
kind: Pod
metadata:
  name: hello-apparmor
  annotations:
    container.apparmor.security.beta.kubernetes.io/hello:localhost/nginx-profile-3
spec:
  containers:
  - name: hello
    image: busybox
    command: [ "sh", "-c", "echo 'Hello AppArmor!' && sleep 1h" ]
```
4. Create after modification


```
$ kubectl apply -f /cks/4/pod1.yaml
```

5. PodSecurityPolicy

switch cluster kubectl config use-context k8s63
context

A PodsecurityPolicy shall prevent the create on of privileged Pods in a specific namespace.Task

Create a new PodSecurityPolicy named prevent-psp-policy , which prevents the creation of privileged Pods.

Create a new ClusterRole named restrict-access-role , which uses the newly created PodSecurityPolicy prevent-psp-policy .

Create a new serviceAccount named pspdenial-sa in the existing namespace development .

Finally, create a new clusterRoleBinding named dany-access-bind , which binds the newly created ClusterRole restrict-access-role to the newly created serviceAccount

Problem solving ideas

PodSecurityPolicy

Keywords: psp policy privileged

0. Switch the group to see if it is enabled

```
$ vi /etc/kubernetes/manifests/kube-apiserver.yaml
```

```
- --enable-admission-plugins=NodeRestriction,PodSecurityPolicy
```

```
$ systemctl restart kubelet
```

1. Copy the psp from the official website, modify the deny privilege

```
$ cat psp.yaml
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: prevent-ppsp-policy
spec:
  privileged: false
  seLinux:
    rule: RunAsAny
  supplementalGroups:
    rule: RunAsAny
  runAsUser:
    rule: RunAsAny
  fsGroup:
    rule: RunAsAny
  volumes:
    - '*'
```

```
$ kubectl create -f psp.yaml
```

2. Create the corresponding clusterrole

```
$ kubectl create clusterrole restrict-access-role --verb=use --resource=podsecuritypolicy --
resource-name=prevent-ppsp-policy
```

3. Create sa to see the corresponding ns

```
$ kubectl create sa psp-denial-sa -n development
```

4. Create a binding relationship

```
$ kubectl create clusterrolebinding dany-access-bind --clusterrole=restrict-access-role --
serviceaccount=development:psp-denial-sa
```

6. Network Policy

switch cluster kubectl config use-context k8s

create a NetworkPolicy named pod-access to restrict access to Pod products-service running in namespace development . only allow the following Pods to connect to Pod products-service :

Pods in the namespace testing

Pods with label environment: staging , in any namespace Make sure to apply the NetworkPolicy.

You can find a skelet on manifest file at/cks/6/p1.yaml

Problem solving ideas

NetworkPolicy

Keywords: NetworkPolicy

1. The host checks the label of the pod

```
$ kubectl get pod -n development --show-labels
```

2. Check the label corresponding to ns, there is no need to set it

```
$ kubectl label ns testing name=testing
```

3. Orchestrate network policy

```
$ cat /cks/6/p1.yaml
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: "pod-access"
```

```
  namespace: "development"
```

```
spec:
```

```
  podSelector:
```

```
    matchLabels:
```

```
      environment: staging
```

```
  policyTypes:
```

```
  - Ingress
```

```
  ingress:
```

```
  - from:
```

```
    - namespaceSelector:
```

```
      matchLabels:
```

```
        name: testing
```

```
  - from:
```

```
    - namespaceSelector:
```

```
      matchLabels:
```

```
        podSelector:
```

```
          matchLabels:
```

```
            environment: staging
```

```
$ kubectl create -f /cks/6/p1.yaml
```

7. dockerfile detection and yaml file problem

switch cluster kubectl config use-context k8s

Task

Analyze and edit the given Dockerfile (based on the ubuntu:16.04 image) /cks/7/Dockerfile fixing two instructions present in the file being prominent security/best-practice issues.

Analyze and edit the given manifest file /cks/7/deployment.yaml

fixing two fields present in the file being prominent security/best-practice issues.

Problem solving ideas

Keywords: Dockerfile issues

1. Pay attention to the number of errors prompted by the dockerfile

Note: USER root

2.yaml problem: pay attention to the api version problem, and the privileged network and mirror version, also depends on the errors mentioned in the title

Case:

Dockerfile

```
# build container stage 1
FROM ubuntu:20.04
ARG DEBIAN_FRONTEND=noninteractive
RUN apt-get update && apt-get install -y golang-go=2:1.13~1ubuntu2
COPY app.go .
RUN pwd
RUN CGO_ENABLED=0 go build app.go
```

```
# app container stage 2
FROM alpine:3.12.0
RUN addgroup -S appgroup && adduser -S appuser -G appgroup -h /home/appuser
RUN rm -rf /bin/*
COPY --from=0 /app /home/appuser/
USER appuser
cmd ["/home/appuser/app"]
```

8. pod security

8. pod security

switch cluster kubectl config use-context k8s
context

It is best-practice to design containers to be stateless and immutable. Task
inspect Pods running in namespace testing and delete any Pod that is either not stateless or not
immutable. use the following strict interpretation of stateless and immutable:
Pods being able to store data inside containers must be treated as not stateless.
You don't have to worry whether data is actually stored inside containers or not already. Pods
being configured to be privileged in any way must be treated as potentially not stateless and not
immutable.

Problem solving ideas

Keywords: stateless immutable

1. get all pods

2. Check if there is a privilege privi*

3. Check if there is volume

4. Delete the privileged network and volume

```
$ kubectl get pod pod1 -n testing -o jsonpath={.spec.volumes} | jq
```

```
$ kubectl get pod sso -n testing -o yaml |grep "privi.*: true"
$ kubectl delete pod xxxxx -n testing
```

9. Create SA

switch cluster kubectl config use-context k8s

context

A Pod fails to run because of an incorrectly specified ServiceAccount.

Task

create a new ServiceAccount named frontend-sa in the existing namespace qa , which must not have access to any secrets. Inspect the Pod named frontend running in the namespace qa . Edit the Pod to use the newly created serviceAccount

Problem solving ideas

Configure Service Accounts for Pods

Keyword: ServiceAccount "must not have access to any secrets"

1. Get the sa template

```
$ kubectl create serviceaccount frontend-sa -n qa --dry-run -o yaml
```

2. Find automatic mounting through official documentation

```
$ k edit pod frontend -n qa
```

apiVersion: v1

kind: Pod

metadata:

creationTimestamp: null

labels:

run: frontend

name: frontend

spec:

serviceAccountName: frontend-sa #Add this line

automountServiceAccountToken: false #Add this line

containers:

- image: nginx

name: frontend

resources: {}

dnsPolicy: ClusterFirst

restartPolicy: Always

status: {}

3. Modify the serviceAccountName in the pod

4. Create pod to delete other sa

10. Trivy detects mirror security

switch cluster kubectl config use-context k8s

Task

Use the Trivy open-source container scanner to detect images with severe vulnerabilities used

by Pods in the namespace yavin . Look for images with High or Critical severity vulnerabilities, and delete the Pods that use those images. Trivy is pre-installed on the cluster's master node only; it is not available on the base system or the worker nodes. You'll have to connect to the cluster's master node to use Trivy

Problem solving ideas

Keywords: Trivy scanner High or Critical

1. Switch the cluster and ssh to the corresponding master
2. get pod scans the corresponding images, no High or Critical
\$ docker run ghcr.io/aquasecurity/trivy:latest image nginx:latest |grep 'High|Critical'
3. Delete the problematic mirror pod
\$ docker rmi <image>

11. Create Secret

switch cluster kubectl config use-context k8s

Task

Retrieve the content of the existing secret named db1-test in the istio-system namespace. store the username field in a file named /cks/11/old-username.txt , and the password field in a file named /cks/11/old-pass.txt. You must create both files; they don't exist yet. Do not use/modify the created files in the following steps, create new temporary files if needed. Create a new secret named test-workflow in the istio-system namespace, with the following content:

username : thanos

password : hahahaha

Finally, create a new Pod that has access to the secret test-workflow via avolume:

pod name dev-pod

namespace istio-system

container name dev-container

image nginx:1.9

volume name dev-volume

mount path /etc/test-secret

Problem solving ideas

Secret

keyword: secret

1. Obtain the username and passwd of db1-test

```
$ kubectl get secrets db1-test -n istio-system -o yaml
```

```
$ echo -n "aGFoYTAwMQ==" | base64 -d > /cks/11/old-pass.txt
```

```
$ echo -n "dG9t" | base64 -d > /cks/11/old-username.txt
```

2. Create a secret named test-workflow

```
$ kubectl create secret generic test-workflow --from-literal=username=thanos --from-literal=password=hahahaha -n istio-system
```

3. More requirements to create pods of secrets

```
$ cat secret-pod.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: dev-pod
```

```
  namespace: istio-system
```

```
spec:
```

```
  containers:
```

```
  - name: dev-container
```

```
    image: nginx:1.9
```

```
    volumeMounts:
```

```
    - name: foo
```

```
      mountPath: "/etc/test-secret"
```

```
      readOnly: true
```

```
  volumes:
```

```
  - name: dev-volume
```

```
    secret:
```

```
      secretName: test-workflow
```

```
k create -f secret-pod.yaml
```

12. kube-bench

```
switch cluster kubectrl config use-context k8s65
```

```
context
```

ACIS Benchmark tool was run against the kubeadm-created cluster and found multiple issues that must be addressed immediately. Task

Fix all issues via configuration and restart the affected components to ensure the new settings take effect. Fix all of the following violations that were found against the API server:

Ensure that the

1.2.7 --authorization-mode FAIL argument is not set to AlwaysAllow

Ensure that the

1.2.8 --authorization-mode FAIL argument includes Node

Ensure that the

1.2.9 --authorization-mode FAIL argument includes RBAC

Ensure that the

1.2.18 --insecure-bind-address FAIL argument is not set

Ensure that the

1.2.19 --insecure-port FAIL argument is set to 0

Fix all of the following violations that were found against the kubelet:

Ensure that the

4.2.1 anonymous-auth FAIL argument is set to false

Ensure that the

4.2.2 --authorization-mode FAIL argument is not set to AlwaysAllow

Use webhook authn/authz

Problem solving ideas

Keywords: look at the entry to determine whether it is a scan

1. Switch the machine to the corresponding ssh to master node

2. kube-bench run Find the corresponding entry, and then fix it

docker run --pid=host -v /etc:/etc:ro -v /var:/var:ro -t aquasec/kube-bench:latest master --version 1.20

There is an ETCD in the exam

Case 1

\$ docker run --pid=host -v /etc:/etc:ro -v /var:/var:ro -t aquasec/kube-bench:latest master --version 1.20

.....

[FAIL] 1.1.12 Ensure that the etcd data directory ownership is set to etcd:etcd (Automated)

.....

\$ cat /etc/kubernetes/kubelet.conf

apiVersion: v1

clusters:

- cluster:

certificate-authority-data:

LS0tLS1CRUdJTiBDRVJUSUZJQ0FURStLS0tCk1JSUM1ekNDQWMrZ0F3SUJBZ0lCQURBTkNa3Foa2lHOXcwQkFRc0ZBREFWTVJNd0VRWURWUWFERXdwcmRXSmwKY201bGRHVnpNQjRYRFRJeE1Ea3hNekE0TVRjd09Wb1hEVE14TURreE1UQTRNVGN3T1Zvd0ZURVRNQkVHQTFVRQpBeE1LYTNWaVpYSnVaWFJsY3pDQ0FTSXdeUUVKS29aSWh2Y05BUUVCQIFBVGdnRVBBRENDQVFvQ2dnRUJBT1BBcnd5ZFJlYlJBWjdGMmpRampHSXFWZlBTZHlVeFMxTElWZTBKaGd0YjFyVXBtdjEydUNtZlVQaEIEUnZ6dktoZnMKexlWnmF2Tm5zZkl2UnpyK3pqMXpDT1gzVFNaYmY0a0NaOE44OEpSSUR0NnBDS0lJU0xiOHVrc3VKtZa5NWVqdGPUvNvVr21CRmVLbGN1ejFHS1FLVEw3aINaNs0TXJNYXlFOUhhkbmJ6dVpNVE42ZlNvRXhGMXhxM29DMGkrZUJCCjNUK1BjMXI5V0NNcndXWEc5VUZmNFo4eFhEaGduL2hESkhKRVJ2eWtsRmpxeGpaRCt4UHILcTg4dk85SytKbVYKbHk2TGw1a21lbDdPdE9ZTURnMWkwUzBJUWZJMgtUaU92d0NsOHM1NVBXUThtTmtZcnJWXXhLTkJBTy94L1FCOAowMFlrNGFmVkrJQlIXZHVIcjFrQ0F3RUFBYU5DTUVBd0RnWURWUjBQQVFIL0JBUURBZ0trTUE4R0ExVWRFd0VCCi93UUZNQU1CQWY4d0hRWURWUjBPQkJZRUZHOu5hRE1RSVQ0c3MwNVpTTGcvV0RibDZ6ZGINQTBHq1Nxr1NjYjMkRFFfQkN3VUFBNlCQVFDaWtKUeD3c0F3YVZ4cWxXazJvR3Vubkc3N1A0ZXFaYlI0d1pMbmbhvdmlRLR1N1Ylg1cgpyMHNrUldHY1RUMTJQQ0xpRWMyaEx2bGp2VC9sUTMvTXV2Nm5iWURQU3g5YjNFb0VGMlI4SHFXTWM1QlhKVS85Ck5wQVhjK20vN01yWkFlcFcx3c3rbmNTVGRIMDFOYIEExQk6ODJrRnpNWU5vSitMSmNFeGx2U2t6ck11V0NXaUEKNnZibGplRnJKQnc1a0UxT3cvR05LOUhuUVFI5OXp1b2U4THJSb2pzUEFUZi92ekFKZExRa2k3MXJpWXRzRkUwNAPZT3JlcUE4Y2ZNeGRuUGNBeG85Z1JWQzBhQzBE2FRcC9aNSStjRTcwVW

```
15dnFQcm44VGJ6MU1uOWt1V2VQTWE5Ck9XNGI3U3RiOVp6NIBCN3pqSzk3dHhzeHRFRW
V1Ky9MeDJXKwotLS0tLUVORCBDRVJUSUZJQ0FURS0tLS0tCg==
```

```
server: https://192.168.211.40:6443
```

```
name: kubernetes
```

```
contexts:
```

```
- context:
```

```
cluster: kubernetes
```

```
user: system:node:master
```

```
name: system:node:master@kubernetes
```

```
current-context: system:node:master@kubernetes
```

```
kind: Config
```

```
preferences: {}
```

```
users:
```

```
- name: system:node:master
```

```
user:
```

```
client-certificate: /var/lib/kubelet/pki/kubelet-client-current.pem
```

```
client-key: /var/lib/kubelet/pki/kubelet-client-current.pem
```

```
$ echo
```

```
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUM1ekNDQWMrZ0F3SUJBZ0lCQURBT
kJna3Foa2IHOXcwQkFRc0ZBREFWTVJNd0VRWURWUUVFERXdwcmRXSmwKY201bGRHVn
pNQjRyRFRJeE1Ea3hNekE0TVRjd09Wb1hEVE14TURreE1UQTRNVGN3T1Zvd0ZURVRNQk
VHQTFVRQpBeE1LYTNWaVpYSnVaWFJsY3pDQ0FTSXdEUVlKS29aSWWhY05BUUVCQlFB
RGdnRVBIBRENDQVFvQ2dnRUJBT1BBcnd5ZFJlYlJBWjdGMmpRampHSXFWZlBTZHVVeFMx
TElwZTBKaGd0YjYjFyVXBtdjEydUNTZlVQaEIEUnZ6dktoZnMKeXlwNmF2Tm5zZkl2UnpyK3pqM
XpDT1gzVFNaYmY0a0NaOE44OEpSSUR0NnBDS0lJU0xiOHVrc3VKZTA5NWVqdGpuVnVvR2
1CRmVLbGN1ejFHS1FLVEw3aWNaNys0TXJNYXlFOUhhbmJ6dVpNVE42ZlNvRXhGMXhxM29
DMGkrZUJCCjNUK1BjMXI5V0NNcndXWEc5VUZmNFo4eFhEaGduL2hESkhKRVJ2eWtsRmpx
eGpaRCt4UHLcTg4dk85SytKbVYKbHk2TGw1a21lbdDpdE9ZTURnMWkwUzBJUWZJMgtUaU
92d0NsOHM1NVBXUThTmtZcnlJWXXhLTkjbTY94L1FCOAowMFIrNGFmVkrRJQlIXZHVIcjFrQ0
F3RUFBYU5DTUVBd0RnWURWUjBQQVFIL0JBUURBZ0trTUE4R0ExVWRFd0VCCi93UUZN
QU1CQWY4d0hRWURWUjBPQkZRUZHOu5hRE1RSVQ0c3MwNVpTTGcvV0RibDZ6ZGINQ
TBHQ1Nxr1NjYjMKRFFFQkN3VUFBNElCQVFDaWtKUeD3c0F3YVZ4cWxXazJvR3Vubkc3N
1A0ZXFaYlI0d1pMbmbhvdmlLR1N1Ylg1cgpyMHNRUldHY1RUMTJQQ0xpRWMyaEx2bGp2VC9s
UTMvTXV2Nm5iWURQU3g5YjNFb0VGMlI4SHFXTWM1QlIhKVS85Ck5wQVhjK20vN01yWkFlc
Fcxc3crbmNTVGRIMDFOYIEExQkY6ODJrRnpNWU5vSitMSmNFeGx2U2t6ck11V0NXaUEKNnZi
bGplRnJKQnc1a0UxT3cvR05LOUhhUVFI5OXp1b2U4THJSb2pzUEFUZi92ekFKZExRa2k3MXJ
pWXRzRkUwNAPZT3JlcUE4Y2ZNeGRuUGNBEG85Z1JWQzBhQzBE2FReC9aNSStjRTcwVW
15dnFQcm44VGJ6MU1uOWt1V2VQTWE5Ck9XNGI3U3RiOVp6NIBCN3pqSzk3dHhzeHRFRW
V1Ky9MeDJXKwotLS0tLUVORCBDRVJUSUZJQ0FURS0tLS0tCg== | base64 -d >
/etc/kubernetes/pki/apiserver-kubelet-ca.crt
```

```
$ cat /etc/kubernetes/pki/apiserver-kubelet-ca.crt
```

```
-----BEGIN CERTIFICATE-----
```

```
MIIC5zCCAC+gAwIBAgIBADANBgkqhkiG9w0BAQsFADAVMRMwEQYDVQQDEwprdWJl
cm5ldGVzMB4XDTE1MDkxMzA4MTcwOV0XDTMxMDkxMTA4MTcwOVowFTETMBEGA1UE
```

```
AxMKa3ViZXJuZXRlczCCASlwDQYJKoZIhvcNAQEBBQADggEPADCCAQoCggEBAOPA
wydYcbRAZ7F2jQjGIlqVfPSdyUxS1LB0e0Jhg1rUpmv12uCmfUPhIDRvzvKhfs
yr06avNnsflvRzr+zj1zCOX3TSZbf4kCZ8N88JRIDt6pCKIISLe8uksuJO095ejv
nVuoGmBFeKlcuz1GKQKTL7jSZ7+4MrMayE9HdnbzuZMTN6fSoExF1xq3oC0i+eBB
3T+Pc1yyWCMrwWVG9UFf4Z8xXDhgn/hDJHJERvyklFjqxjZD+xPyKq88vO9K+JmV
ly6LI5kme17OtOYMDg1i0S0IQfI0kTiOvwCl8s55PWQ8mNkYrylYxKNBAO/x/QB8
00Yk4afVDIBYWduHr1kCAwEAANCMEEAwDgYDVR0PAQH/BAQDAgKkMA8GA1UdEwEB
/wQFMAMBAf8wHQYDVR0OBBYEFG9NaDMQIT4ss05ZSLg/WDbI6zdiMA0GCSqGSIb3
DQEBEwUAA4IBAQCikJPGwsAwaVxqlWk2oGunnG77P4eqZbYtwZLnhoViKGSuX5r
r0skRWGcTT12PCLiEc2hLvljvT/IQ3/Muv6nbYDPSx9b3EoEF2YxHqWMc5BXJU/9
NpAXc+m/7MrZAepW1sw+ncSTdH01NbQ1BBz82kFzMYNoJ+LJcExlvSkzrMuWCWiA
6vbljHFrJBw5kE1Ow/GNK9HTTR99zuoe8LrRojSPATf/vzAJdLQki71riYtsFE04
YOreqA8cfMxdnPcAxo9gRVC0aC0DwaQx/Z5+cE70UmyvqPrn8Tbz1Mn9kuWePMA9
OW4b7Stb9Zz6PB7zjK97txsxtEEeu+/Lx2W+
-----END CERTIFICATE-----
```

```
$ vim /etc/kubernetes/manifests/kube-apiserver.yaml
```

```
.....
```

```
--kubecertificates-authority=/etc/kubernetes/pki/apiserver-kubecertificates-ca.crt
```

```
.....
```

```
$ kubectl get pods -n kube-system | grep kube-apiserver
```

```
$ docker run --pid=host -v /etc:/etc:ro -v /var:/var:ro -t aquasec/kube-bench:latest master --
version 1.20
```

13. gVisor

Change the cluster kubectl config use-context k8s67

context

This cluster uses containerd as CRI runtime. Containerd's default runtime handler is runc .

Containerd has been prepared to support an additional runtime handler ,runsc (gVisor). Task:

Create a RuntimeClass named untrusted using the prepared runtime handler named runsc .

Update all Pods in the namespace client to run on gvisor, unless they are already running on a non-default runtime handler. You can find a skeleton manifest file at /cks/13/rc.yaml

Problem solving ideas

RuntimeClass

Keywords: gVisor

1. Switch the cluster and create a runtimeclass with the official website documentation

```
$ vim rc.yaml
```

```
apiVersion: node.k8s.io/v1beta1
```

```
kind: RuntimeClass
```

```
metadata:
```



```
name: untrusted
handler: runsc
```

```
$ k -f rc.yaml create
```

2. Another topic requires creating a pod to use this runtime

```
$ k edit pod mypod -n client
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  namespace: client
spec:
  runtimeClassName: untrusted
.....
```

14. Audit

switch cluster kubectl config use-context k8s

Task

Enable audit logs in the cluster. To do so, enable the log backend, and ensure that:

logs are stored at /var/log/kubernetes/audit-logs.txt

log files are retained for 5 days at maximum, a number of 10 auditlog files are retained

A basic policy is provided at /etc/kubernetes/logpolicy/sample-policy.yaml . it only specifies what not to log. The base policy is located on the cluster's master node. Edit and extend the basic policy to log:

namespaces changes at RequestResponse level

the request body of pods changes in the namespace front-apps

configMap and secret changes in all namespaces at the Metadata level

Also, add a catch-all rule to log all other requests at the Metadata level. Don't forget to apply

Problem solving ideas

audit

keyword: policy

1. Switch the cluster to log in to the master, then create a directory, modify yaml, and enable auditing

```
$ mkdir /var/log/kubernetes/
```

```
$ mkdir /etc/kubernetes/logpolicy/
```

```
$ cat /etc/kubernetes/logpolicy/sample-policy.yaml
```

```
$ cat policy.yaml
```

```
apiVersion: audit.k8s.io/v1
```

```
kind: Policy
```

```
omitStages:
```

```
- "RequestReceived"
```

```
rules:
```

- level: RequestResponse
 resources:
 - group: ""
 resources: ["namespaces"]
- level: Request
 resources:
 - group: "" # core API group
 resources: ["pods"]
 namespaces: ["front-apps"]
- level: Metadata
 resources:
 - group: ""
 resources: ["secrets", "configmaps"]
- level: Metadata
 omitStages:
 - "RequestReceived"

2. More official website documents to modify the corresponding strategy

```
$ vim /etc/kubernetes/manifests/kube-apiserver.yaml
```

- --audit-policy-file=/etc/kubernetes/logpolicy/sample-policy.yaml # add
- --audit-log-path=/var/log/kubernetes/audit-logs.txt # add
- --audit-log-maxage=5 # add
- --audit-log-maxbackup=10

.....

- mountPath: /etc/kubernetes/logpolicy # add
- name: audit # add
- hostNetwork: true
- priorityClassName: system-node-critical
- volumes:
 - hostPath: # add
 path: /etc/kubernetes/logpolicy # add
 type: DirectoryOrCreate # add
 name: audit # add

3. Restart kubelet

```
$ systemctl restart kubelet
```

```
$ k get pods -n kube-system | grep api
```

```
$ cat /var/log/kubernetes/audit-logs.txt
```

15. Default Network Policy

```
switch cluster kubectl config use-context k8s
context
```

A default-deny NetworkPolicy avoids to accident all y expose a Pod in a namespace that doesn't have any other NetworkPolicy defined. Task

Create a new default-deny NetworkPolicy named denynetwork in the namespace development for all traffic of type Ingress . The new NetworkPolicy must deny all Ingress traffic in the namespace development . Apply the newly created default-deny NetworkPolicy to all Pods running in namespace development . You can find a skeleton manifest file

Problem solving ideas
NetworkPolicy

Keywords: NetworkPolicy defined

1. Observe clearly whether all conditions are rejected by default or other conditions, and more topics require official documents to write yaml

```
$ cat denynetwork.yaml
```

```
apiVersion: networking.k8s.io/v1
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: denynetwork
```

```
  namespace: development
```

```
spec:
```

```
  podSelector: {}
```

```
  policyTypes:
```

```
  - Ingress
```

```
$ k create -f denynetwork.yaml
```

16. Falco detection output log format

```
$ ssh node1
```

```
$ systemctl stop falco
```

```
$ falco
```

```
$ cd /etc/falco/
```

```
$ ls
```

```
falco_rules.local.yaml falco_rules.yaml falco.yaml k8s_audit_rules.yaml rules.available rules.d
```

```
$ rep -r "A shell was spawned in a container with an attached terminal" *
```

```
falco_rules.yaml: A shell was spawned in a container with an attached terminal
```

```
(user=%user.name user_loginuid=%user.loginuid %container.info)
```

```
#update configuration
```

```
root@node1:/etc/falco# cat falco_rules.local.yaml
- rule: Terminal shell in container
  desc: A shell was used as the entrypoint/exec point into a container with an attached terminal.
  condition: >
    spawned_process and container
    and shell_procs and proc.tty != 0
    and container_entrypoint
    and not user_expected_terminal_shell_in_container_conditions
  output: >
    %evt.time,%user.name,%container.name,%container.id
    shell=%proc.name parent=%proc.pname cmdline=%proc.cmdline terminal=%proc.tty
    container_id=%container.id image=%container.image.repository)
  priority: WARNING
  tags: [container, shell, mitre_execution]
```

```
$ falco
Mon May 24 00:07:13 2021: Falco version 0.28.1 (driver version
5c0b863ddade7a45568c0ac97d037422c9efb750)
Mon May 24 00:07:13 2021: Falco initialized with configuration file /etc/falco/falco.yaml
Mon May 24 00:07:13 2021: Loading rules from file /etc/falco/falco_rules.yaml:
Mon May 24 00:07:13 2021: Loading rules from file /etc/falco/falco_rules.local.yaml:
#Configuration takes effect
Mon May 24 00:07:13 2021: Loading rules from file /etc/falco/k8s_audit_rules.yaml:
Mon May 24 00:07:14 2021: Starting internal webserver, listening on port 8765
00:07:30.297671117: Warning Shell history had been deleted or renamed (user=root
user_loginuid=-1 type=openat command=bash fd.name=/root/.bash_history
name=/root/.bash_history path=<NA> oldpath =<NA> k8s_apache_apache_default_3ece2efb-
fe49-4111-899f-10d38a61bab6_0 (id=84dd6fe8a9ad))
```

```
format change
00:07:33.763063865: Warning
00:07:33.763063865,root,k8s_apache_apache_default_3ece2efb-fe49-4111-899f-
10d38a61bab6_0,84dd6fe8a9ad shell=bash parent=runc cmdline=bash terminal=34816
container_id=84addd6fe8a9)
```

<https://blog.csdn.net/xixihahalelehehe/article/details/122525427>